

Logistic Regression Project

In this project we will be working with the UCI adult dataset. We will be attempting to predict if people in the data set belong in a certain class by salary, either making $\leq 50k$ or $> 50k$ per year.

Typically most of your time is spent cleaning data, not running the few lines of code that build your model, this project will try to reflect that by showing different issues that may arise when cleaning data.

Get the Data

Read in the `adult_sal.csv` file and set it to a data frame called `adult`.

In [1]:

```
setwd("C:/Users/ADMIN/Desktop/DataScience/R")
adult<-read.csv("adult_sal.csv",sep = ";",header = TRUE)

#names(adult)<-c("age","workclass","fnlwgt","education","education-num","marital-status","occupation","relations
hip","race","sex","capital-gain","capital-loss","hours-per-week","native-country")
```

Check the head of `adult`

In [2]:

```
head(adult)
Out[2]:
```

	X	age	type_employer	fnlwgt	education	education_num	marital	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_week	country	income
1	1	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	$\leq 50K$
2	2	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Executive-manager	Husband	White	Male	0	0	13	United-States	$\leq 50K$

			inc				ous e	ial							tes	
3	3	38	Pri vate	21 564 6	HS -gra d	9	Di vorc ed	Ha ndle rs-cl ean ers	No t-in- fami ly	W hite	Ma le	0	0	40	Un ited -Sta tes	<= 50K
4	4	53	Pri vate	23 472 1	11 th	7	Ma rried -civ -sp ous e	Ha ndle rs-cl ean ers	Hu sba nd	Bl ack	Ma le	0	0	40	Un ited -Sta tes	<= 50K
5	5	28	Pri vate	33 840 9	Ba chel ors	13	Ma rried -civ -sp ous e	Pr of-s peci alty	Wi fe	Bl ack	Fe mal e	0	0	40	Cu ba	<= 50K
6	6	37	Pri vate	28 458 2	Ma ster s	14	Ma rried -civ -sp ous e	Ex ec- man ager ial	Wi fe	W hite	Fe mal e	0	0	40	Un ited -Sta tes	<= 50K

You should notice the index has been repeated. Drop this column.

In [3]:

```
library(dplyr)
adult <- select(adult,-X)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Check the head, str, and summary of the data now.

In []:

head(adult)

	age	type_employer	fnlwgt	education	education_num	marital	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_week	country	income
1	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<= 50K
2	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<= 50K
3	38	Private	215646	HS-grad	9	Divorced	Handler-scleanser	Not-in-family	White	Male	0	0	40	United-States	<= 50K
4	53	Private	234721	11th	7	Married-civ-spouse	Handler-scleanser	Husband	Black	Male	0	0	40	United-States	<= 50K
5	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<= 50K

						se									
6	37	Pri vate	28 4582	Ma sters	14	Ma rried -civ- spou se	Ex ec- man ageri al	Wif e	Wh ite	Fe male	0	0	40	Uni ted- Stat es	<= 50K

In []:

str(adult)

```
## 'data.frame':      48842 obs. of  15 variables:

##  $ age                : int   25 38 28 44 18 34 29 63 24 55 ...

##  $ workclass           : Factor w/  9 levels "?","Federal-gov",...: 5 5 3 5 1 5 1 7 5 5 ...

##  $ fnlwgt              : int   226802 89814 336951 160323 103497 198693 227026 104626 369667
104996 ...

##  $ education           : Factor w/ 16 levels "10th","11th",...: 2 12 8 16 16 1 12 15 16 6 ...

##  $ educational.num     : int    7  9 12 10 10 6  9 15 10 4 ...

##  $ marital.status      : Factor w/  7 levels "Divorced","Married-AF-spouse",...: 5 3 3 3 5 5 5 3 5
3 ...

##  $ occupation          : Factor w/ 15 levels "?","Adm-clerical",...: 8 6 12 8 1 9 1 11 9 4 ...

##  $ relationship        : Factor w/  6 levels "Husband","Not-in-family",...: 4 1 1 1 4 2 5 1 5 1
...

##  $ race                : Factor w/  5 levels "Amer-Indian-Eskimo",...: 3 5 5 3 5 5 3 5 5 5 ...

##  $ gender              : Factor w/  2 levels "Female","Male": 2 2 2 2 1 2 2 2 1 2 ...

##  $ capital.gain         : int    0  0  0 7688 0  0  0 3103 0  0 ...

##  $ capital.loss         : int    0  0  0  0  0  0  0  0  0  0 ...

##  $ hours.per.week       : int   40 50 40 40 30 30 40 32 40 10 ...

##  $ native.country       : Factor w/ 42 levels "?","Cambodia",...: 40 40 40 40 40 40 40 40 40 40
...

##  $ income               : Factor w/  2 levels "<=50K", ">50K": 1 1 2 2 1 1 1 2 1 1 ...
```

In []:

```
> summary(adult)

      X1              X2              X3              X4              X5
Min.   :17.00  Length:32561  Min.    : 12285  Length:32561  Min.    : 1.00
```

1st Qu.:28.00	Class :character	1st Qu.: 117827	Class :character	1st Qu.: 9.00
Median :37.00	Mode :character	Median : 178356	Mode :character	Median :10.00
Mean :38.58		Mean : 189778		Mean :10.08
3rd Qu.:48.00		3rd Qu.: 237051		3rd Qu.:12.00
Max. :90.00		Max. :1484705		Max. :16.00

X6	X7	X8	X9	X10
Length:32561	Length:32561	Length:32561	Length:32561	Length:32561
Class :character	Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character	Mode :character

X11	X12	X13	X14	X15
Min. : 0	Min. : 0.0	Min. : 1.00	Length:32561	Length:32561
1st Qu.: 0	1st Qu.: 0.0	1st Qu.:40.00	Class :character	Class :character
Median : 0	Median : 0.0	Median :40.00	Mode :character	Mode :character
Mean : 1078	Mean : 87.3	Mean :40.44		
3rd Qu.: 0	3rd Qu.: 0.0	3rd Qu.:45.00		
Max. :99999	Max. :4356.0	Max. :99.00		

Data Cleaning

Notice that we have a lot of columns that are cateogrical factors, however a lot of these columns have too many factors than may be necessary. In this data cleaning section we'll try to clean these columns up by reducing the number of factors.

type_employer column

Use `table()` to check out the frequency of the `type_employer` column.

In [7]:

```
> table(adult$X2)
```

Out[7]:

```
      ?      Federal-gov      Local-gov      Never-worked
1836      960      2093           7
Private  Self-emp-inc Self-emp-not-inc   State-gov
22696      1116      2541      1298
Without-pay
14
```

How many Null values are there for `type_employer`? What are the two smallest groups?

In [8]:

```
> sum(is.na(adult$X2))
[1] 0
```

```
x10<-adult%>%group_by(X10)
```

```
x10
```

```
# A tibble: 32,561 x 15
# Groups:   X10 [2]
      X1 X2      X3 X4      X5 X6      X7      X8      X9      X10      X11      X12      X13 X14      X15
  <int> <chr>    <int> <chr> <int> <chr>    <chr>    <chr> <chr> <chr> <int> <int> <int> <chr> <chr>
1    39 State--  77516 Bache~  13 Never-- Adm-c~ Not-- White Male    2174    0    40 Unit~ <=50K
2    50 Self-e~  83311 Bache~  13 Marrie~ Exec-- Husb~ White Male      0    0    13 Unit~ <=50K
3    38 Private 215646 HS-gr~    9 Divorc~ Handl~ Not-- White Male      0    0    40 Unit~ <=50K
4    53 Private 234721 11th      7 Marrie~ Handl~ Husb~ Black Male      0    0    40 Unit~ <=50K
5    28 Private 338409 Bache~  13 Marrie~ Prof-- Wife  Black Fema~    0    0    40 Cuba  <=50K
6    37 Private 284582 Maste~  14 Marrie~ Exec-- Wife  White Fema~    0    0    40 Unit~ <=50K
7    49 Private 160187 9th      5 Marrie~ Other~ Not-- Black Fema~    0    0    16 Jama~ <=50K
8    52 Self-e~ 209642 HS-gr~    9 Marrie~ Exec-- Husb~ White Male      0    0    45 Unit~ >50K
9    31 Private  45781 Maste~  14 Never-- Prof-- Not-- White Fema~ 14084    0    50 Unit~ >50K
10   42 Private 159449 Bache~  13 Marrie~ Exec-- Husb~ White Male    5178    0    40 Unit~ >50K
# ... with 32,551 more rows
```

```
x15<-adult%>%group_by(X15)
x15
```

```
# A tibble: 32,561 x 15
# Groups:   X15 [2]
      X1 X2      X3 X4      X5 X6      X7      X8      X9      X10      X11      X12      X13 X14      X15
  <int> <chr>    <int> <chr> <int> <chr>    <chr>    <chr> <chr> <chr> <int> <int> <int> <chr> <chr>
1    39 State--  77516 Bache~  13 Never-- Adm-c~ Not-- White Male    2174    0    40 Unit~ <=50K
2    50 Self-e~  83311 Bache~  13 Marrie~ Exec-- Husb~ White Male      0    0    13 Unit~ <=50K
3    38 Private 215646 HS-gr~    9 Divorc~ Handl~ Not-- White Male      0    0    40 Unit~ <=50K
4    53 Private 234721 11th      7 Marrie~ Handl~ Husb~ Black Male      0    0    40 Unit~ <=50K
5    28 Private 338409 Bache~  13 Marrie~ Prof-- Wife  Black Fema~    0    0    40 Cuba  <=50K
6    37 Private 284582 Maste~  14 Marrie~ Exec-- Wife  White Fema~    0    0    40 Unit~ <=50K
7    49 Private 160187 9th      5 Marrie~ Other~ Not-- Black Fema~    0    0    16 Jama~ <=50K
8    52 Self-e~ 209642 HS-gr~    9 Marrie~ Exec-- Husb~ White Male      0    0    45 Unit~ >50K
9    31 Private  45781 Maste~  14 Never-- Prof-- Not-- White Fema~ 14084    0    50 Unit~ >50K
10   42 Private 159449 Bache~  13 Marrie~ Exec-- Husb~ White Male    5178    0    40 Unit~ >50K
# ... with 32,551 more rows
```

Combine these two smallest groups into a single group called "Unemployed". There are lots of ways to do this, so feel free to get creative. Hint: It may be helpful to convert these objects into character data types (as.character()) and then use supply with a custom function)

```
adult$X2 <- as.character(adult$X2)
```

In [9]:

In [10]:

```
adult$X2[adult$X2 == "Without-pay" |  
        adult$X2 == "Never-worked"] <- "Unemployed"
```

In [11]:

```
table(adult$X2)
```

Out[11]:

?	Federal-gov	Local-gov	Private
1836	960	2093	22696
Self-emp-inc	Self-emp-not-inc	State-gov	Unemployed
1116	2541	1298	21

What other columns are suitable for combining? Combine State and Local gov jobs into a category called SL-gov and combine self-employed jobs into a category called self-emp.

In [12]:

```
adult$X2[adult$X2 == "State-gov" |  
        adult$X2 == "Local-gov"] <- "SL-gov"
```

In [13]:

```
adult$X2[adult$X2 == "Self-emp-inc" |  
        adult$X2 == "Self-emp-not-inc"] <- "Self-emp"
```

In [14]:

```
table(adult$X2)
```

Out[14]:

?	Federal-gov	Private	self-emp	SL-gov	Unemployed
1836	960	22696	3657	3391	21

Marital Column

Use table() to look at the marital column

In [15]:

```
table(adult$X6)
Notex6:marital column
```

Out[15]:

Divorced	Married-AF-spouse	Married-civ-spouse
4443	23	14976
Married-spouse-absent	Never-married	Separated
418	10683	1025
Widowed		
993		

Reduce this to three groups:

- Married
- Not-Married
- Never-Married

In [16]:

```
adult$X6 <- as.character(adult$X6)

adult$X6[adult$X6 == "Married-AF-spouse" |
          adult$X6 == "Married-civ-spouse" |
          adult$X6 == "Married-spouse-absent"] <- "Married"

adult$X6[adult$X6 == "Divorced" |
          adult$X6 == "Separated" |
          adult$X6 == "Widowed"] <- "Not-Married"
```


In [17]:

```
table (adult$X6)
```

Out[

Married	Never-married	Not-Married
15417	10683	6461

Country Column

Check the country column using table()

In [18]:

```
table(adult$x14)  
x14 means country
```

Out[18]:

?	Cambodia
583	19
Canada	China
121	75
Columbia	Cuba
59	95
Dominican-Republic	Ecuador
70	28
El-Salvador	England
106	90
France	Germany
29	137
Greece	Guatemala
29	64
Haiti	Holand-Netherlands
44	1
Honduras	Hong
13	20
Hungary	India
13	100
Iran	Ireland
43	24
Italy	Jamaica
73	81
Japan	Laos

62	18	
Mexico	Nicaragua	
643	34	
Outlying-US(Guam-USVI-etc)		Peru
14	31	
Philippines	Poland	
198	60	
Portugal	Puerto-Rico	
37	114	
Scotland	South	
12	80	
Taiwan	Thailand	
51	18	
Trinidad&Tobago	United-States	
19	29170	
Vietnam	Yugoslavia	
67	16	

Group these countries together however you see fit. You have flexibility here because there is no right/wrong way to do this, possibly group by continents. You should be able to reduce the number of groups here significantly though.

In [19]:

```
adult$X14 <- as.character(adult$X14)
```

Out[19]:

1. '?'
2. 'Cambodia'
3. 'Canada'
4. 'China'
5. 'Columbia'
6. 'Cuba'
7. 'Dominican-Republic'
8. 'Ecuador'
9. 'El-Salvador'
10. 'England'
11. 'France'
12. 'Germany'
13. 'Greece'
14. 'Guatemala'
15. 'Haiti'

16. 'Holand-Netherlands'
17. 'Honduras'
18. 'Hong'
19. 'Hungary'
20. 'India'
21. 'Iran'
22. 'Ireland'
23. 'Italy'
24. 'Jamaica'
25. 'Japan'
26. 'Laos'
27. 'Mexico'
28. 'Nicaragua'
29. 'Outlying-US(Guam-USVI-etc)'
30. 'Peru'
31. 'Philippines'
32. 'Poland'
33. 'Portugal'
34. 'Puerto-Rico'
35. 'Scotland'
36. 'South'
37. 'Taiwan'
38. 'Thailand'
39. 'Trinidad&Tobago'
40. 'United-States'
41. 'Vietnam'
42. 'Yugoslavia'

In [20]:

```
north.america <- c("Canada", "Cuba", "Dominican-Republic", "El-Salvador", "Guatemala",  
                  "Haiti", "Honduras", "Jamaica", "Mexico", "Nicaragua",  
                  "Outlying-US(Guam-USVI-etc)", "Puerto-Rico", "Trinidad&Tobago",  
                  "United-States")  
  
asia <- c("Cambodia", "China", "Hong", "India", "Iran", "Japan", "Laos",  
          "Philippines", "Taiwan", "Thailand", "Vietnam")  
  
south.america <- c("Columbia", "Ecuador", "Peru")  
  
europe <- c("England", "France", "Germany", "Greece", "Holand-Netherlands",
```

```
"Hungary", "Ireland", "Italy", "Poland", "Portugal", "Scotland",  
"Yugoslavia")
```

```
other <- c("South", "?")
```

In [21]:

```
adult$X14[adult$X14 %in% north.america] <- "North America"  
adult$X14[adult$X14 %in% asia] <- "Asia"  
adult$X14[adult$X14 %in% south.america] <- "South America"  
adult$X14[adult$X14 %in% europe] <- "Europe"  
adult$X14[adult$X14 %in% other] <- "Other"
```

In [22]:

```
table(adult$X14)
```

Use table() to confirm the groupings

In [23]:

```
table(adult$X14)
```

Out[23]

Asia	Europe	Latin.and.South.America
671	521	1301
North.America	Other	
29405	663	

Check the str() of adult again. Make sure any of the columns we changed have factor levels with factor()

In []:

```
adult$X14 <- as.factor(adult$X14)  
adult$X6 <- as.factor(adult$X6)  
adult$X2 <- as.factor(adult$X2)
```

In [25]:

```
str(adult)
```

In []:

We could still play around with education and occupation to try to reduce the number of factors for those columns, but let's go ahead and move on to dealing with the missing data. Feel free to group those columns as well and see how they effect your model.

Missing Data ¶

Notice how we have data that is missing.

Amelia ¶

Install and load the Amelia package.

In [27]:

```
#install.packages('Amelia',repos = 'http://cran.us.r-project.org')  
library(Amelia)
```

Loading required package: Rcpp

```
##  
## Amelia II: Multiple Imputation  
## (Version 1.7.4, built: 2015-12-05)  
## Copyright (C) 2005-2016 James Honaker, Gary King and Matthew Blackwell  
## Refer to http://gking.harvard.edu/amelia/ for more information  
##
```

Convert any cell with a '?' or a ' ' value to a NA value. Hint: `is.na()` may be useful here or you can also use brackets with a conditional statement. Refer to the solutions if you can't figure this step out.

In [28]:

Using `table()` on a column with NA values should now not display those NA values, instead you'll just see 0 for ?. Optional: Refactor these columns (may take awhile). For example:

In [29]:

```
adult[adult == "?"] <- NA  
  
table(adult$X2)
```

Out[29]:

SL-gov	self-emp	Private	Federal-gov	?	Unemployed
3391	3657	22696	960	0	21

In [30]:


Play around with the `missmap` function from the `Amelia` package. Can you figure out what its doing and how to use it?

In [31]:

```
missmap(adult,y.at=c(1),y.labels = c(""),col=c('red','black'))
```


Missingness Map

Missing Observed



You should have noticed that using `missmap(adult)` is basically a heatmap pointing out missing values (NA). This gives you a quick glance at how much data is missing, in this case, not a whole lot (relatively speaking). You probably also noticed that there is a bunch of y labels, get rid of them by running the command below. What is `col=c('yellow','black')` doing?

```
missmap(adult, y.at = 1, y.labels = "", col = c("yellow", "black"))
```



Use `na.omit()` to omit NA data from the adult data frame. Note, it really depends on the situation and your data to judge

whether or not this is a good decision. You shouldn't always just drop NA values.

In [33]:

```
adult <- na.omit(adult)
```

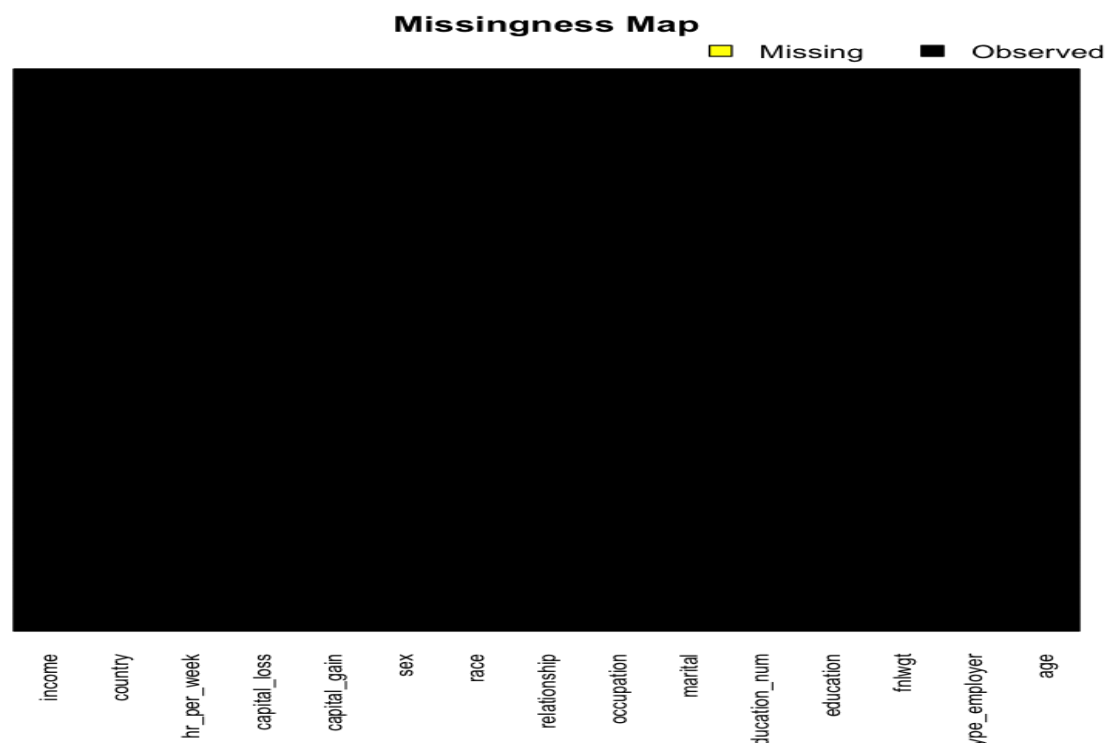
```
# May take awhile
```

```
#str(adult)
```

Use `missmap()` to check that all the NA values were in fact dropped.

In [34]:

```
library(Amelia)
missmap(adult, legend = True, col=c('yellow','black'))
```



Although we've cleaned the data, we still have explored it using visualization.

Check the str() of the data.

In [35]:

Str(adult)

```
'data.frame': 30718 obs. of 15 variables:
 $ age      : int  39 50 38 53 28 37 49 52 31 42 ...
 $ type_employer: Factor w/ 5 levels "SL-gov","self-emp",...: 1 2 3 3 3 3 2 3 3 ...
 $ fnlwgt    : int  77516 83311 215646 234721 338409 284582 160187 209642 45781 159449 ...
 $ education  : Factor w/ 16 levels "10th","11th",...: 10 10 12 2 10 13 7 12 13 10 ...
 $ education_num: int  13 13 9 7 13 14 5 9 14 13 ...
 $ marital    : Factor w/ 3 levels "Never-married",...: 1 2 3 2 2 2 2 2 1 2 ...
 $ occupation  : Factor w/ 14 levels "Adm-clerical",...: 1 2 3 3 4 2 5 2 4 2 ...
 $ relationship : Factor w/ 6 levels "Husband","Not-in-family",...: 2 1 2 1 6 6 2 1 2 1 ...
 $ race        : Factor w/ 5 levels "Amer-Indian-Eskimo",...: 5 5 5 3 3 5 3 5 5 5 ...
 $ sex         : Factor w/ 2 levels "Female","Male": 2 2 2 2 1 1 1 2 1 2 ...
 $ capital_gain : int  2174 0 0 0 0 0 0 0 14084 5178 ...
 $ capital_loss : int  0 0 0 0 0 0 0 0 0 0 ...
 $ hr_per_week  : int  40 13 40 40 40 40 16 45 50 40 ...
 $ country      : Factor w/ 5 levels "North.America",...: 1 1 1 1 2 1 2 1 1 1 ...
 $ income       : Factor w/ 2 levels "<=50K",">50K": 1 1 1 1 1 1 1 2 2 2 ...
 - attr(*, "na.action")=Class 'omit' Named int [1:1843] 28 62 70 78 107 129 150 155 161 188 ...
 .. ..- attr(*, "names")= chr [1:1843] "28" "62" "70" "78" ...
```


Use ggplot2 to create a histogram of ages, colored by income.

In [36]:

library(ggplot2)

In [37]:

```
ggplot(data = adult,aes(x=age,fill=income))+geom_histogram(col='black')
```




Plot a histogram of hours worked per week

In [38]:

```
ggplot(data = adult,aes(x=hr_per_week))+geom_histogram(col='black')
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



Rename the country column to region column to better reflect the factor levels.

In [52]:

```
names(adult)[14]<-"region"
```


In []:

```
Region<-factor(region)  
levels(Region)<-c("North.America","Latin.South.and.America","Asia","other","Europe")
```

Create a barplot of region with the fill color defined by income class. Optional: Figure out how rotate the x axis text for readability

In [54]:

```
ggplot(data = adult,aes(x=region,fill=income))+geom_bar(col='black')
```



Now it's time to build a model to classify people into two groups: Above or Below 50k in Salary.

Logistic Regression¶

Refer to the Lecture or ISLR if you are fuzzy on any of this.

Logistic Regression is a type of classification model. In classification models, we attempt to predict the outcome of categorical dependent variables, using one or more independent variables. The independent variables can be either categorical or numerical.

Logistic regression is based on the logistic function, which always takes values between 0 and 1. Replacing the dependent variable of the logistic function with a linear combination of dependent variables we intend to use for regression, we arrive at the formula for logistic regression.

Take a quick look at the head() of adult to make sure we have a good overview before going into building the model!

In [55]:

Head(adult)

Out[55]:

	age	type_employer	fnlwgt	education	education_num	marital	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_week	region	income
1	39	SL-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	North.America	<= 50K
2	50	self-emp	83311	Bachelors	13	Married	Exec-managerial	Husband	White	Male	0	0	13	North.America	<= 50K
3	38	Private	215646	HS-grad	9	Not-Married	Handlers-cleaners	Not-in-family	White	Male	0	0	40	North.America	<= 50K

4	53	Private	234721	11th	7	Married	Handler-s-cleaners	Husband	Black	Male	0	0	40	North.America	<= 50K
5	28	Private	338409	Bachelors	13	Married	Prof-specialty	Wife	Black	Female	0	0	40	Latin.and.South.America	<= 50K
6	37	Private	284582	Masters	14	Married	Exec-managerial	Wife	White	Female	0	0	40	North.America	<= 50K

Train Test Split

Split the data into a train and test set using the `caTools` library as done in previous lectures. Reference previous solutions notebooks if you need a refresher.

In [59]:

```
library(caTools)
set.seed(123)
split = sample.split(adult$income, SplitRatio = 0.75)
training_set = subset(adult, split == TRUE)
test_set = subset(adult, split == FALSE)
```

Training the Model

Explore the `glm()` function with `help(glm)`. Read through the documentation.

In [60]:

help(glm)

Out[60]:

glm {stats}	R Documentation
-------------	-----------------

Fitting Generalized Linear Models

Description

glm is used to fit generalized linear models, specified by giving a symbolic description of the linear predictor and a description of the error distribution.

Usage

```
glm(formula, family = gaussian, data, weights, subset,
    na.action, start = NULL, etastart, mustart, offset,
    control = list(...), model = TRUE, method = "glm.fit",
    x = FALSE, y = TRUE, contrasts = NULL, ...)
```

```
glm.fit(x, y, weights = rep(1, nobs),
    start = NULL, etastart = NULL, mustart = NULL,
    offset = rep(0, nobs), family = gaussian(),
    control = list(), intercept = TRUE)
```

```
## S3 method for class 'glm'
weights(object, type = c("prior", "working"), ...)
```

Arguments

formula	an object of class "formula" (or one that can be coerced to that class): a symbolic description of the model to be fitted. The details of model specification are given under ‘Details’.
family	a description of the error distribution and link function to be used in the model. For glm this can be a character string naming a family function, a family function or the result of a call to a family function. For glm.fit only the third option is supported. (See family for details of family functions.)

data	an optional data frame, list or environment (or object coercible by <code>as.data.frame</code> to a data frame) containing the variables in the model. If not found in data, the variables are taken from <code>environment(formula)</code> , typically the environment from which <code>glm</code> is called.
weights	an optional vector of ‘prior weights’ to be used in the fitting process. Should be <code>NULL</code> or a numeric vector.
subset	an optional vector specifying a subset of observations to be used in the fitting process.
na.action	a function which indicates what should happen when the data contain NAs. The default is set by the <code>na.action</code> setting of options, and is <code>na.fail</code> if that is unset. The ‘factory-fresh’ default is <code>na.omit</code> . Another possible value is <code>NULL</code> , no action. Value <code>na.exclude</code> can be useful.
start	starting values for the parameters in the linear predictor.
etastart	starting values for the linear predictor.
mustart	starting values for the vector of means.
offset	this can be used to specify an <i>a priori</i> known component to be included in the linear predictor during fitting. This should be <code>NULL</code> or a numeric vector of length equal to the number of cases. One or more offset terms can be included in the formula instead or as well, and if more than one is specified their sum is used. See <code>model.offset</code> .
control	a list of parameters for controlling the fitting process. For <code>glm.fit</code> this is passed to <code>glm.control</code> .
model	a logical value indicating whether <i>model frame</i> should be included as a component of the returned value.
method	the method to be used in fitting the model. The default method <code>"glm.fit"</code> uses iteratively reweighted least squares (IWLS): the alternative <code>"model.frame"</code> returns the model frame and does no fitting. User-supplied fitting functions can be supplied

	either as a function or a character string naming a function, with a function which takes the same arguments as <code>glm.fit</code> . If specified as a character string it is looked up from within the stats namespace.
x, y	For <code>glm</code> : logical values indicating whether the response vector and model matrix used in the fitting process should be returned as components of the returned value. For <code>glm.fit</code> : x is a design matrix of dimension $n \times p$, and y is a vector of observations of length n.
contrasts	an optional list. See the <code>contrasts.arg</code> of <code>model.matrix.default</code> .
intercept	logical. Should an intercept be included in the <i>null</i> model?
object	an object inheriting from class "glm".
type	character, partial matching allowed. Type of weights to extract from the fitted model object. Can be abbreviated.
...	For <code>glm</code> : arguments to be used to form the default control argument if it is not supplied directly. For weights: further arguments passed to or from other methods.

Details

A typical predictor has the form `response ~ terms` where `response` is the (numeric) response vector and `terms` is a series of terms which specifies a linear predictor for response. For binomial and quasibinomial families the response can also be specified as a factor (when the first level denotes failure and all others success) or as a two-column matrix with the columns giving the numbers of successes and failures. A terms specification of the form `first + second` indicates all the terms in `first` together with all the terms in `second` with any duplicates removed.

A specification of the form `first:second` indicates the set of terms obtained by taking the interactions of all terms in `first` with all terms in `second`. The specification `first*second` indicates the *cross* of `first` and `second`. This is the same as `first + second + first:second`.

The terms in the formula will be re-ordered so that main effects come first, followed by the interactions, all second-order, all third-order and so on: to avoid this pass a terms object as the formula.

Non-NULL weights can be used to indicate that different observations have different dispersions (with the values in `weights` being inversely proportional to the dispersions); or equivalently, when the elements of `weights` are positive integers w_i , that each response y_i is the mean of w_i unit-weight observations. For a binomial GLM prior weights are used to give the number of trials when the response is the proportion of successes: they would rarely be used for a Poisson GLM.

`glm.fit` is the workhorse function: it is not normally called directly but can be more efficient where the response vector, design

matrix and family have already been calculated.

If more than one of `etastart`, `start` and `mustart` is specified, the first in the list will be used. It is often advisable to supply starting values for a quasi family, and also for families with unusual links such as `gaussian("log")`.

All of `weights`, `subset`, `offset`, `etastart` and `mustart` are evaluated in the same way as variables in formula, that is first in data and then in the environment of formula.

For the background to warning messages about ‘fitted probabilities numerically 0 or 1 occurred’ for binomial GLMs, see Venables & Ripley (2002, pp. 197–8).

Value

`glm` returns an object of class inheriting from `"glm"` which inherits from the class `"lm"`. See later in this section. If a non-standard method is used, the object will also inherit from the class (if any) returned by that function.

The function `summary` (i.e., `summary.glm`) can be used to obtain or print a summary of the results and the function `anova` (i.e., `anova.glm`) to produce an analysis of variance table.

The generic accessor functions `coefficients`, `effects`, `fitted.values` and `residuals` can be used to extract various useful features of the value returned by `glm`.

`weights` extracts a vector of weights, one for each case in the fit (after subsetting and `na.action`).

An object of class `"glm"` is a list containing at least the following components:

coefficients	a named vector of coefficients
residuals	the <i>working</i> residuals, that is the residuals in the final iteration of the IWLS fit. Since cases with zero weights are omitted, their working residuals are NA.
fitted.values	the fitted mean values, obtained by transforming the linear predictors by the inverse of the link function.
rank	the numeric rank of the fitted linear model.
family	the family object used.
linear.predictors	the linear fit on link scale.
deviance	up to a constant, minus twice the maximized log-likelihood. Where sensible, the constant is chosen so that a saturated model has deviance zero.
aic	A version of Akaike's <i>An Information Criterion</i> , minus twice the maximized log-likelihood plus twice the number of parameters, computed by the <code>aic</code> component of the family. For binomial and Poisson families the dispersion is fixed at one and the number of parameters is the number of

	coefficients. For gaussian, Gamma and inverse gaussian families the dispersion is estimated from the residual deviance, and the number of parameters is the number of coefficients plus one. For a gaussian family the MLE of the dispersion is used so this is a valid value of AIC, but for Gamma and inverse gaussian families it is not. For families fitted by quasi-likelihood the value is NA.
null.deviance	The deviance for the null model, comparable with deviance. The null model will include the offset, and an intercept if there is one in the model. Note that this will be incorrect if the link function depends on the data other than through the fitted mean: specify a zero offset to force a correct calculation.
iter	the number of iterations of IWLS used.
weights	the <i>working</i> weights, that is the weights in the final iteration of the IWLS fit.
prior.weights	the weights initially supplied, a vector of 1s if none were.
df.residual	the residual degrees of freedom.
df.null	the residual degrees of freedom for the null model.
y	if requested (the default) the y vector used. (It is a vector even for a binomial model.)
x	if requested, the model matrix.
model	if requested (the default), the model frame.
converged	logical. Was the IWLS algorithm judged to have converged?
boundary	logical. Is the fitted value on the boundary of the attainable values?
call	the matched call.
formula	the formula supplied.

terms	the terms object used.
data	the data argument.
offset	the offset vector used.
control	the value of the control argument used.
method	the name of the fitter function used, currently always "glm.fit".
contrasts	(where relevant) the contrasts used.
xlevels	(where relevant) a record of the levels of the factors used in fitting.
na.action	(where relevant) information returned by model.frame on the special handling of NAs.

In addition, non-empty fits will have components `qr`, `R` and effects relating to the final weighted linear fit.

Objects of class "glm" are normally of class `c("glm", "lm")`, that is inherit from class "lm", and well-designed methods for class "lm" will be applied to the weighted linear model at the final iteration of IWLS. However, care is needed, as extractor functions for class "glm" such as `residuals` and `weights` do **not** just pick out the component of the fit with the same name.

If a binomial glm model was specified by giving a two-column response, the `weights` returned by `prior.weights` are the total numbers of cases (factored by the supplied case weights) and the component `y` of the result is the proportion of successes.

Fitting functions

The argument `method` serves two purposes. One is to allow the model frame to be recreated with no fitting. The other is to allow the default fitting function `glm.fit` to be replaced by a function which takes the same arguments and uses a different fitting algorithm. If `glm.fit` is supplied as a character string it is used to search for a function of that name, starting in the `stats` namespace.

The class of the object return by the fitter (if any) will be prepended to the class returned by `glm`.

Author(s)

The original `R` implementation of `glm` was written by Simon Davies working for Ross Ihaka at the University of Auckland, but has since been extensively re-written by members of the R Core team.

The design was inspired by the `S` function of the same name described in Hastie & Pregibon (1992).

References

- Dobson, A. J. (1990) *An Introduction to Generalized Linear Models*. London: Chapman and Hall.
- Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.
- McCullagh P. and Nelder, J. A. (1989) *Generalized Linear Models*. London: Chapman and Hall.
- Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. New York: Springer.

See Also

`anova.glm`, `summary.glm`, etc. for `glm` methods, and the generic functions `anova`, `summary`, `effects`, `fitted.values`, and `residuals`.
`lm` for non-generalized *linear* models (which SAS calls GLMs, for ‘general’ linear models).
`loglin` and `loglm` (package [MASS](#)) for fitting log-linear models (which binomial and Poisson GLMs are) to contingency tables.
`bigglm` in package [biglm](#) for an alternative way to fit GLMs to large datasets (especially those with many cases).
`esoph`, `infert` and `predict.glm` have examples of fitting binomial glms.

Examples

```
## Dobson (1990) Page 93: Randomized Controlled Trial :
```

```
counts <- c(18,17,15,20,10,20,25,13,12)
outcome <- gl(3,1,9)
treatment <- gl(3,3)
print(d.AD <- data.frame(treatment, outcome, counts))
glm.D93 <- glm(counts ~ outcome + treatment, family = poisson())
anova(glm.D93)
summary(glm.D93)
```

```
## an example with offsets from Venables & Ripley (2002, p.189)
```

```
utils::data(anorexia, package = "MASS")
```

```
anorex.1 <- glm(Postwt ~ Prewt + Treat + offset(Prewt),
               family = gaussian, data = anorexia)
summary(anorex.1)
```

```
# A Gamma example, from McCullagh & Nelder (1989, pp. 300-2)
```

```
clotting <- data.frame(
  u = c(5,10,15,20,30,40,60,80,100),
  lot1 = c(118,58,42,35,27,25,21,19,18),
  lot2 = c(69,35,26,21,18,16,13,12,12))
summary(glm(lot1 ~ log(u), data = clotting, family = Gamma))
summary(glm(lot2 ~ log(u), data = clotting, family = Gamma))
```

```
## Not run:
```

```
## for an example of the use of a terms object as a formula
demo(glm.vr)
```

```
## End(Not run)
```

Use all the features to train a glm() model on the training data set, pass the argument family=binomial(logit) into the glm function.

In [61]:

```
model = glm(formula = income ~ .,  
            family = binomial(logit),  
            data = training_set)
```

Warning message:
: glm.fit: fitted probabilities numerically 0 or 1 occurred

If you get a warning, this just means that the model may have guessed the probability of a class with a 0% or 100% chance of occuring.

Check the model summary

In [62]:

```
Summary(model)
```

Out[62]:

Call:
glm(formula = income ~ ., family = binomial(logit), data = train)

Deviance Residuals:

Min	1Q	Median	3Q	Max
-5.1163	-0.5172	-0.1965	0.0000	3.6235

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-7.364e+00	4.245e-01	-17.346	< 2e-16 ***
age	2.534e-02	2.007e-03	12.627	< 2e-16 ***
type_employerself-emp	7.501e-03	8.999e-02	0.083	0.933571
type_employerPrivate	2.371e-01	7.321e-02	3.239	0.001198 **
type_employerFederal-gov	6.835e-01	1.266e-01	5.399	6.71e-08 ***
type_employerUnemployed	-1.346e+01	3.688e+02	-0.036	0.970888
fnlwgt	5.424e-07	2.085e-07	2.601	0.009291 **
education11th	2.094e-01	2.570e-01	0.814	0.415384
education12th	3.925e-01	3.410e-01	1.151	0.249612
education1st-4th	-4.590e-01	6.067e-01	-0.757	0.449323
education5th-6th	-8.009e-02	3.980e-01	-0.201	0.840503
education7th-8th	-4.991e-01	2.880e-01	-1.733	0.083096 .
education9th	-1.229e-02	3.191e-01	-0.038	0.969292

educationAssoc-acdm	1.250e+00	2.165e-01	5.775	7.70e-09	***
educationAssoc-voc	1.452e+00	2.084e-01	6.970	3.17e-12	***
educationBachelors	2.003e+00	1.938e-01	10.337	< 2e-16	***
educationDoctorate	2.874e+00	2.636e-01	10.902	< 2e-16	***
educationHS-grad	8.359e-01	1.888e-01	4.426	9.58e-06	***
educationMasters	2.347e+00	2.063e-01	11.374	< 2e-16	***
educationPreschool	-1.879e+01	1.645e+02	-0.114	0.909053	
educationProf-school	2.797e+00	2.468e-01	11.337	< 2e-16	***
educationSome-college	1.203e+00	1.915e-01	6.283	3.33e-10	***
education_num	NA	NA	NA	NA	
maritalMarried	1.280e+00	1.943e-01	6.588	4.45e-11	***
maritalNot-Married	5.435e-01	9.953e-02	5.460	4.75e-08	***
occupationExec-managerial	7.689e-01	9.095e-02	8.453	< 2e-16	***
occupationHandlers-cleaners	-7.944e-01	1.726e-01	-4.603	4.17e-06	***
occupationProf-specialty	4.957e-01	9.626e-02	5.149	2.62e-07	***
occupationOther-service	-8.248e-01	1.386e-01	-5.952	2.65e-09	***
occupationSales	2.896e-01	9.749e-02	2.971	0.002972	**
occupationCraft-repair	4.151e-02	9.483e-02	0.438	0.661616	
occupationTransport-moving	-1.114e-01	1.189e-01	-0.937	0.348928	
occupationFarming-fishing	-1.120e+00	1.619e-01	-6.920	4.52e-12	***
occupationMachine-op-inspct	-2.194e-01	1.203e-01	-1.824	0.068080	.
occupationTech-support	6.829e-01	1.325e-01	5.153	2.56e-07	***
occupationProtective-serv	6.029e-01	1.491e-01	4.044	5.24e-05	***
occupationArmed-Forces	-6.252e-01	1.844e+00	-0.339	0.734504	
occupationPriv-house-serv	-3.600e+00	1.938e+00	-1.858	0.063179	.
relationshipNot-in-family	-8.661e-01	1.907e-01	-4.541	5.60e-06	***
relationshipOther-relative	-1.086e+00	2.546e-01	-4.268	1.97e-05	***
relationshipOwn-child	-1.797e+00	2.357e-01	-7.625	2.45e-14	***
relationshipUnmarried	-1.031e+00	2.154e-01	-4.784	1.72e-06	***
relationshipWife	1.476e+00	1.235e-01	11.949	< 2e-16	***
raceAsian-Pac-Islander	6.073e-01	3.206e-01	1.894	0.058243	.
raceBlack	4.528e-01	2.847e-01	1.590	0.111800	
raceOther	4.135e-02	4.217e-01	0.098	0.921902	
raceWhite	6.595e-01	2.711e-01	2.432	0.014997	*
sexMale	8.855e-01	9.378e-02	9.442	< 2e-16	***
capital_gain	3.192e-04	1.273e-05	25.076	< 2e-16	***
capital_loss	6.549e-04	4.561e-05	14.358	< 2e-16	***
hr_per_week	2.906e-02	1.987e-03	14.623	< 2e-16	***
regionLatin.and.South.America	-5.925e-01	1.595e-01	-3.714	0.000204	***
regionAsia	-6.475e-02	2.044e-01	-0.317	0.751446	
regionOther	-4.300e-01	1.651e-01	-2.604	0.009206	**
regionEurope	4.404e-02	1.552e-01	0.284	0.776660	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 24138 on 21502 degrees of freedom
 Residual deviance: 14004 on 21449 degrees of freedom
 AIC: 14112

Number of Fisher Scoring iterations: 14

We have still a lot of features! Some important, some not so much. R comes with an *awesome* function called `step()`. The `step()` function iteratively tries to remove predictor variables from the model in an attempt to delete variables that do not significantly add to the fit. How does it do this? It uses [AIC](#). Read the wikipedia page for AIC if you want to further understand this, you can also check out `help(step)`. This level of statistics is outside the scope of this project assignment so let's keep moving along

In [63]:

```
help(step)
```

Out[63]:

step {stats}	R Documentation
--------------	-----------------

Choose a model by AIC in a Stepwise Algorithm

Description

Select a formula-based model by AIC.

Usage

```
step(object, scope, scale = 0,  
      direction = c("both", "backward", "forward"),  
      trace = 1, keep = NULL, steps = 1000, k = 2, ...)
```

Arguments

object	an object representing a model of an appropriate class (mainly "lm" and "glm"). This is used as the initial model in the stepwise search.
scope	defines the range of models examined in the stepwise search. This should be either a single formula, or a list containing components upper and lower, both formulae. See the details for how to specify the formulae and how they are used.
scale	used in the definition of the AIC statistic for selecting the models, currently only for lm, aov and glm models. The default value, 0, indicates the scale should be estimated: see <code>extractAIC</code> .

direction	the mode of stepwise search, can be one of "both", "backward", or "forward", with a default of "both". If the scope argument is missing the default for direction is "backward". Values can be abbreviated.
trace	if positive, information is printed during the running of step. Larger values may give more detailed information.
keep	a filter function whose input is a fitted model object and the associated AIC statistic, and whose output is arbitrary. Typically keep will select a subset of the components of the object and return them. The default is not to keep anything.
steps	the maximum number of steps to be considered. The default is 1000 (essentially as many as required). It is typically used to stop the process early.
k	the multiple of the number of degrees of freedom used for the penalty. Only $k = 2$ gives the genuine AIC: $k = \log(n)$ is sometimes referred to as BIC or SBC.
...	any additional arguments to <code>extractAIC</code> .

Details

step uses `add1` and `drop1` repeatedly; it will work for any method for which they work, and that is determined by having a valid method for `extractAIC`. When the additive constant can be chosen so that AIC is equal to Mallows' C_p , this is done and the tables are labelled appropriately.

The set of models searched is determined by the scope argument. The right-hand-side of its lower component is always included in the model, and right-hand-side of the model is included in the upper component. If scope is a single formula, it specifies the upper component, and the lower model is empty. If scope is missing, the initial model is used as the upper model.

Models specified by scope can be templates to update object as used by `update.formula`. So using `.` in a scope formula means ‘what is already there’, with `.^2` indicating all interactions of existing terms.

There is a potential problem in using glm fits with a variable scale, as in that case the deviance is not simply related to the maximized log-likelihood. The "glm" method for function `extractAIC` makes the appropriate adjustment for a gaussian family, but may need to be amended for other cases. (The binomial and poisson families have fixed scale by default and do not correspond to a particular maximum-likelihood problem for variable scale.)

Value

the stepwise-selected model is returned, with up to two additional components. There is an "anova" component corresponding to the steps taken in the search, as well as a "keep" component if the `keep=` argument was supplied in the call. The "Resid. Dev"

column of the analysis of deviance table refers to a constant minus twice the maximized log likelihood: it will be a deviance only in cases where a saturated model is well-defined (thus excluding `lm`, `aov` and `survreg` fits, for example).

Warning

The model fitting must apply the models to the same dataset. This may be a problem if there are missing values and `R`'s default of `na.action = na.omit` is used. We suggest you remove the missing values first.

Calls to the function `nobs` are used to check that the number of observations involved in the fitting process remains unchanged.

Note

This function differs considerably from the function in `S`, which uses a number of approximations and does not in general compute the correct AIC.

This is a minimal implementation. Use `stepAIC` in package [MASS](#) for a wider range of object classes.

Author(s)

B. D. Ripley: `step` is a slightly simplified version of `stepAIC` in package [MASS](#) (Venables & Ripley, 2002 and earlier editions).

The idea of a step function follows that described in Hastie & Pregibon (1992); but the implementation in `R` is more general.

References

Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J.

Hastie, Wadsworth & Brooks/Cole.

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. New York: Springer (4th ed).

See Also

`stepAIC` in [MASS](#), `add1`, `drop1`

Examples

```
## following on from example(lm)
```

```
step(lm.D9)
```

```
summary(lm1 <- lm(Fertility ~ ., data = swiss))
```

```
slm1 <- step(lm1)
```

```
summary(slm1)
```

```
slm1$anova
```

[Package *stats* version 3.2.2]

Use `new.model <- step(your.model.name)` to use the `step()` function to create a new model.

```
New.model<-step(model)
```

Start: AIC=14112.05

```
income ~ age + type_employer + fnlwgt + education + education_num +  
marital + occupation + relationship + race + sex + capital_gain +  
capital_loss + hr_per_week + region
```

Warning message:

[illegible]

Step: AIC=14112.05

```
income ~ age + type_employer + fnlwgt + education + marital +  
  occupation + relationship + race + sex + capital_gain + capital_loss +  
  hr_per_week + region
```

Warning message:

[illegible]

Df Deviance AIC

```
<none> 14004 14112
- fnlwgt 1 14011 14117
- race 4 14019 14119
- region 4 14026 14126
- type_employer 4 14050 14150
- marital 2 14060 14164
- sex 1 14097 14203
- age 1 14165 14271
- capital_loss 1 14217 14323
- hr_per_week 1 14222 14328
- relationship 5 14288 14386
- occupation 13 14444 14526
- education 15 14718 14796
- capital_gain 1 15248 15354
```

You should get a bunch of messages informing you of the process. Check the new.model by using summary()

In [66]:

Summary(new.model)

Out[66]:

Call:

```
glm(formula = income ~ age + type_employer + fnlwgt + education +
    marital + occupation + relationship + race + sex + capital_gain +
    capital_loss + hr_per_week + region, family = binomial(logit),
    data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-5.1163	-0.5172	-0.1965	0.0000	3.6235

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-7.364e+00	4.245e-01	-17.346	< 2e-16 ***
age	2.534e-02	2.007e-03	12.627	< 2e-16 ***
type_employerself-emp	7.501e-03	8.999e-02	0.083	0.933571
type_employerPrivate	2.371e-01	7.321e-02	3.239	0.001198 **
type_employerFederal-gov	6.835e-01	1.266e-01	5.399	6.71e-08 ***
type_employerUnemployed	-1.346e+01	3.688e+02	-0.036	0.970888
fnlwgt	5.424e-07	2.085e-07	2.601	0.009291 **
education11th	2.094e-01	2.570e-01	0.814	0.415384
education12th	3.925e-01	3.410e-01	1.151	0.249612
education1st-4th	-4.590e-01	6.067e-01	-0.757	0.449323
education5th-6th	-8.009e-02	3.980e-01	-0.201	0.840503
education7th-8th	-4.991e-01	2.880e-01	-1.733	0.083096 .
education9th	-1.229e-02	3.191e-01	-0.038	0.969292
educationAssoc-acdm	1.250e+00	2.165e-01	5.775	7.70e-09 ***
educationAssoc-voc	1.452e+00	2.084e-01	6.970	3.17e-12 ***

educationBachelors	2.003e+00	1.938e-01	10.337	< 2e-16 ***
educationDoctorate	2.874e+00	2.636e-01	10.902	< 2e-16 ***
educationHS-grad	8.359e-01	1.888e-01	4.426	9.58e-06 ***
educationMasters	2.347e+00	2.063e-01	11.374	< 2e-16 ***
educationPreschool	-1.879e+01	1.645e+02	-0.114	0.909053
educationProf-school	2.797e+00	2.468e-01	11.337	< 2e-16 ***
educationSome-college	1.203e+00	1.915e-01	6.283	3.33e-10 ***
maritalMarried	1.280e+00	1.943e-01	6.588	4.45e-11 ***
maritalNot-Married	5.435e-01	9.953e-02	5.460	4.75e-08 ***
occupationExec-managerial	7.689e-01	9.095e-02	8.453	< 2e-16 ***
occupationHandlers-cleaners	-7.944e-01	1.726e-01	-4.603	4.17e-06 ***
occupationProf-specialty	4.957e-01	9.626e-02	5.149	2.62e-07 ***
occupationOther-service	-8.248e-01	1.386e-01	-5.952	2.65e-09 ***
occupationSales	2.896e-01	9.749e-02	2.971	0.002972 **
occupationCraft-repair	4.151e-02	9.483e-02	0.438	0.661616
occupationTransport-moving	-1.114e-01	1.189e-01	-0.937	0.348928
occupationFarming-fishing	-1.120e+00	1.619e-01	-6.920	4.52e-12 ***
occupationMachine-op-inspct	-2.194e-01	1.203e-01	-1.824	0.068080 .
occupationTech-support	6.829e-01	1.325e-01	5.153	2.56e-07 ***
occupationProtective-serv	6.029e-01	1.491e-01	4.044	5.24e-05 ***
occupationArmed-Forces	-6.252e-01	1.844e+00	-0.339	0.734504
occupationPriv-house-serv	-3.600e+00	1.938e+00	-1.858	0.063179 .
relationshipNot-in-family	-8.661e-01	1.907e-01	-4.541	5.60e-06 ***
relationshipOther-relative	-1.086e+00	2.546e-01	-4.268	1.97e-05 ***
relationshipOwn-child	-1.797e+00	2.357e-01	-7.625	2.45e-14 ***
relationshipUnmarried	-1.031e+00	2.154e-01	-4.784	1.72e-06 ***
relationshipWife	1.476e+00	1.235e-01	11.949	< 2e-16 ***
raceAsian-Pac-Islander	6.073e-01	3.206e-01	1.894	0.058243 .
raceBlack	4.528e-01	2.847e-01	1.590	0.111800
raceOther	4.135e-02	4.217e-01	0.098	0.921902
raceWhite	6.595e-01	2.711e-01	2.432	0.014997 *
sexMale	8.855e-01	9.378e-02	9.442	< 2e-16 ***
capital_gain	3.192e-04	1.273e-05	25.076	< 2e-16 ***
capital_loss	6.549e-04	4.561e-05	14.358	< 2e-16 ***
hr_per_week	2.906e-02	1.987e-03	14.623	< 2e-16 ***
regionLatin.and.South.America	-5.925e-01	1.595e-01	-3.714	0.000204 ***
regionAsia	-6.475e-02	2.044e-01	-0.317	0.751446
regionOther	-4.300e-01	1.651e-01	-2.604	0.009206 **
regionEurope	4.404e-02	1.552e-01	0.284	0.776660

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 24138 on 21502 degrees of freedom
Residual deviance: 14004 on 21449 degrees of freedom
AIC: 14112

Number of Fisher Scoring iterations: 14

You should have noticed that the `step()` function kept all the features used previously! While we used the AIC criteria to compare models, there are other criteria we could have used. If you want you can try reading about the variable inflation factor (VIF) and `vif()` function to explore other options for comparison criteria. In the meantime let's continue on and see how well our model performed against the test set.

[Review what a confusion matrix is on wikipedia.](#)

Create a confusion matrix using the `predict` function with `type='response'` as an argument inside of that function.

In [74]:

```
y_pred = predict(model, type = 'response', newdata = test_set)
cm = table(test_set, y_pred)
```

Warning message:

In predict.lm(object, newdata, se.fit, scale = 1, type = ifelse(type == : prediction from a rank-deficient fit may be misleading

Out[74]:

```
      FALSE TRUE
<=50K  6372  548
>50K   872 1423
```

You'll notice we have a rank deficient fit. Find out more about what issues this may cause by reading this [stackexchange post](#).

What was the accuracy of our model?

In [76]:

```
(6372+1423)/(6372+548+872+1423)
```

Out[76]:

```
0.845903418339664
```

Calculate other measures of performance like, recall or precision.

In [77]:

$(6372)/(6372+872)$

Out[77]:

0.972832369942197

In [78]:

$(6372)/(6372+548)$

Out[78]:

0.929320817228051

In your opinion, how good was this model? What other context would you like to know before answering that question?

No right/wrong answers here, just want you to think about accuracy, precision, and recall. You would like to know the costs associated with each.

Great Job! 